

Principles of HDFS & MapReduce2:

Name: อวยศัย ภิรมย์รัตน์

Tel. : 086-813-5354


e-mail : p.Auoychai@gmail.com

Big Data

Wrap up

HDFS Architecture :

- **What is HDFS** , Java based distributed file system and allow to store large data across multiple nodes in a Hadoop Cluster.



Distributed Storage:

MasterNode:

NameNode , ResourceManager

SlaveNode:

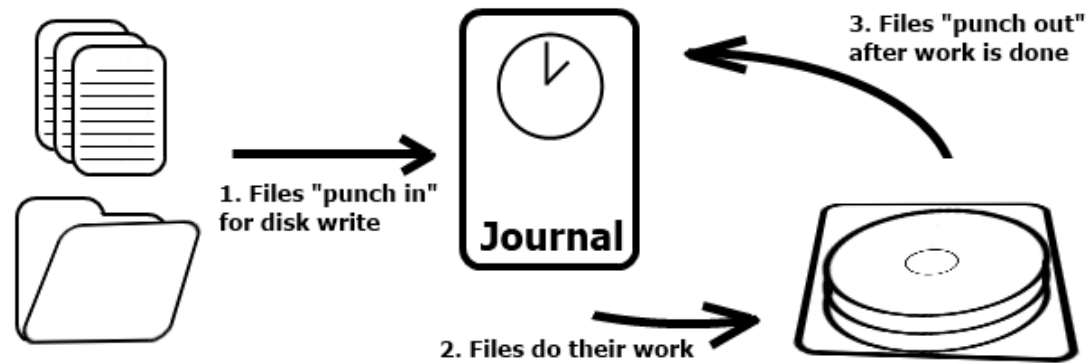
DataNode , NodeManager

Can't use Linux command to see it

The screenshot shows a terminal window with the command `ish@inode ~ $ ls -l` and its output. The output lists various files and directories with their permissions, owner, group, size, and date. The files include `Calibre Library`, `colors-conkyl-gs.zip`, `colors-conkyl-u.zip`, `conky_bu`, `conky-nadia-lb.zip`, `demorec`, `Desktop`, `Documents`, `Downloads`, `Music`, `pen_backup`, `Pictures`, `Public`, `Shared`, `Templates`, `to_print`, `Videos`, and `VirtualBox VMs`.



HDFS Write-Read Process :

What the HDFS difference than traditional File System



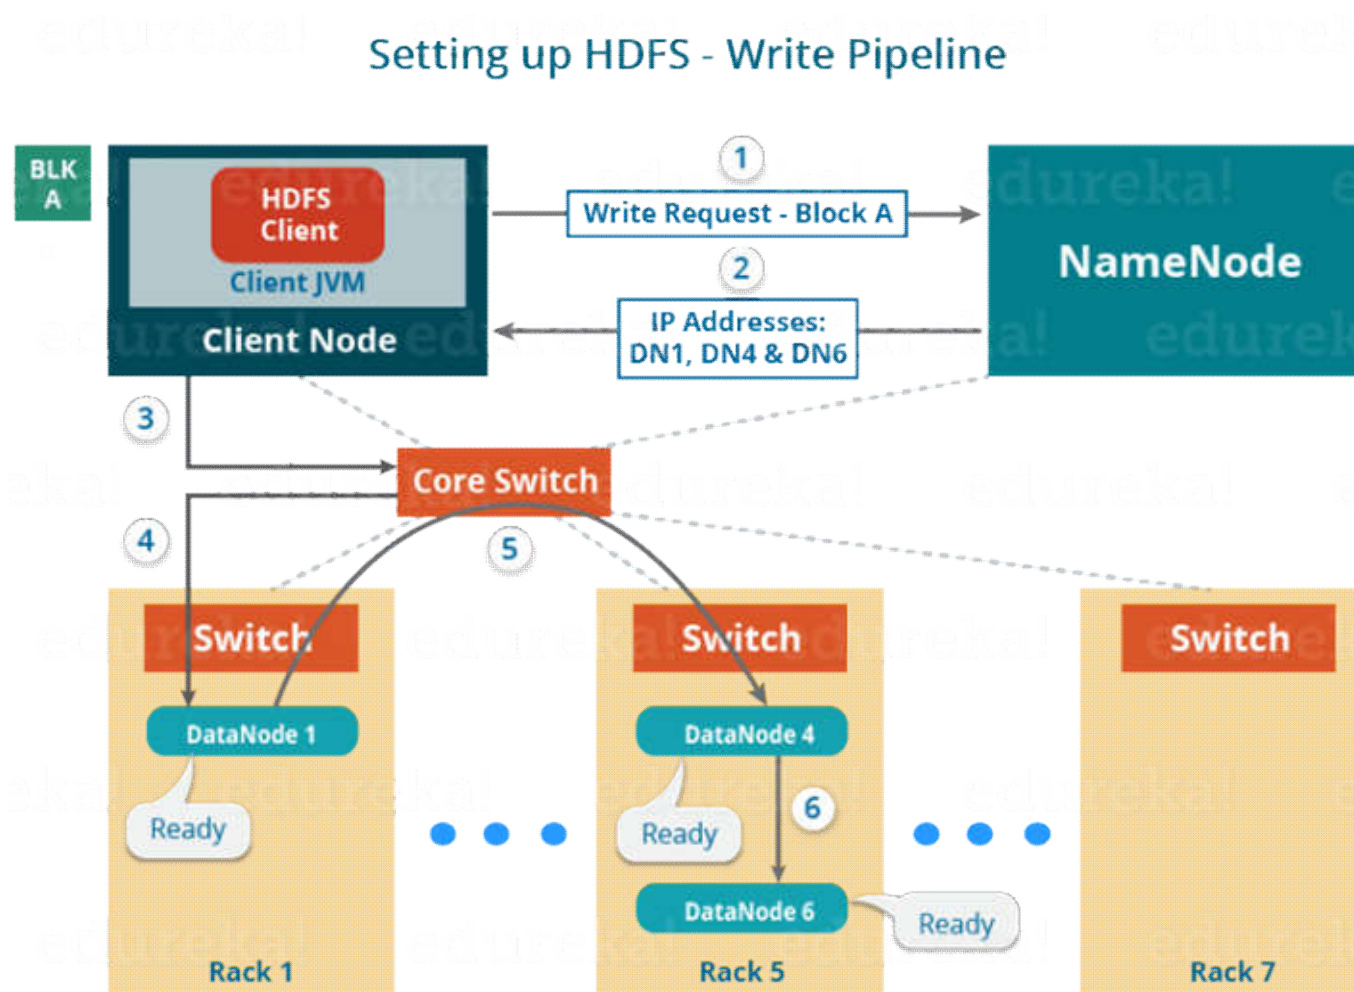
HDFS Write Process :

1. Set up of Pipeline
2. Data streaming and replication
3. Shutdown of Pipeline (Acknowledgement stage)



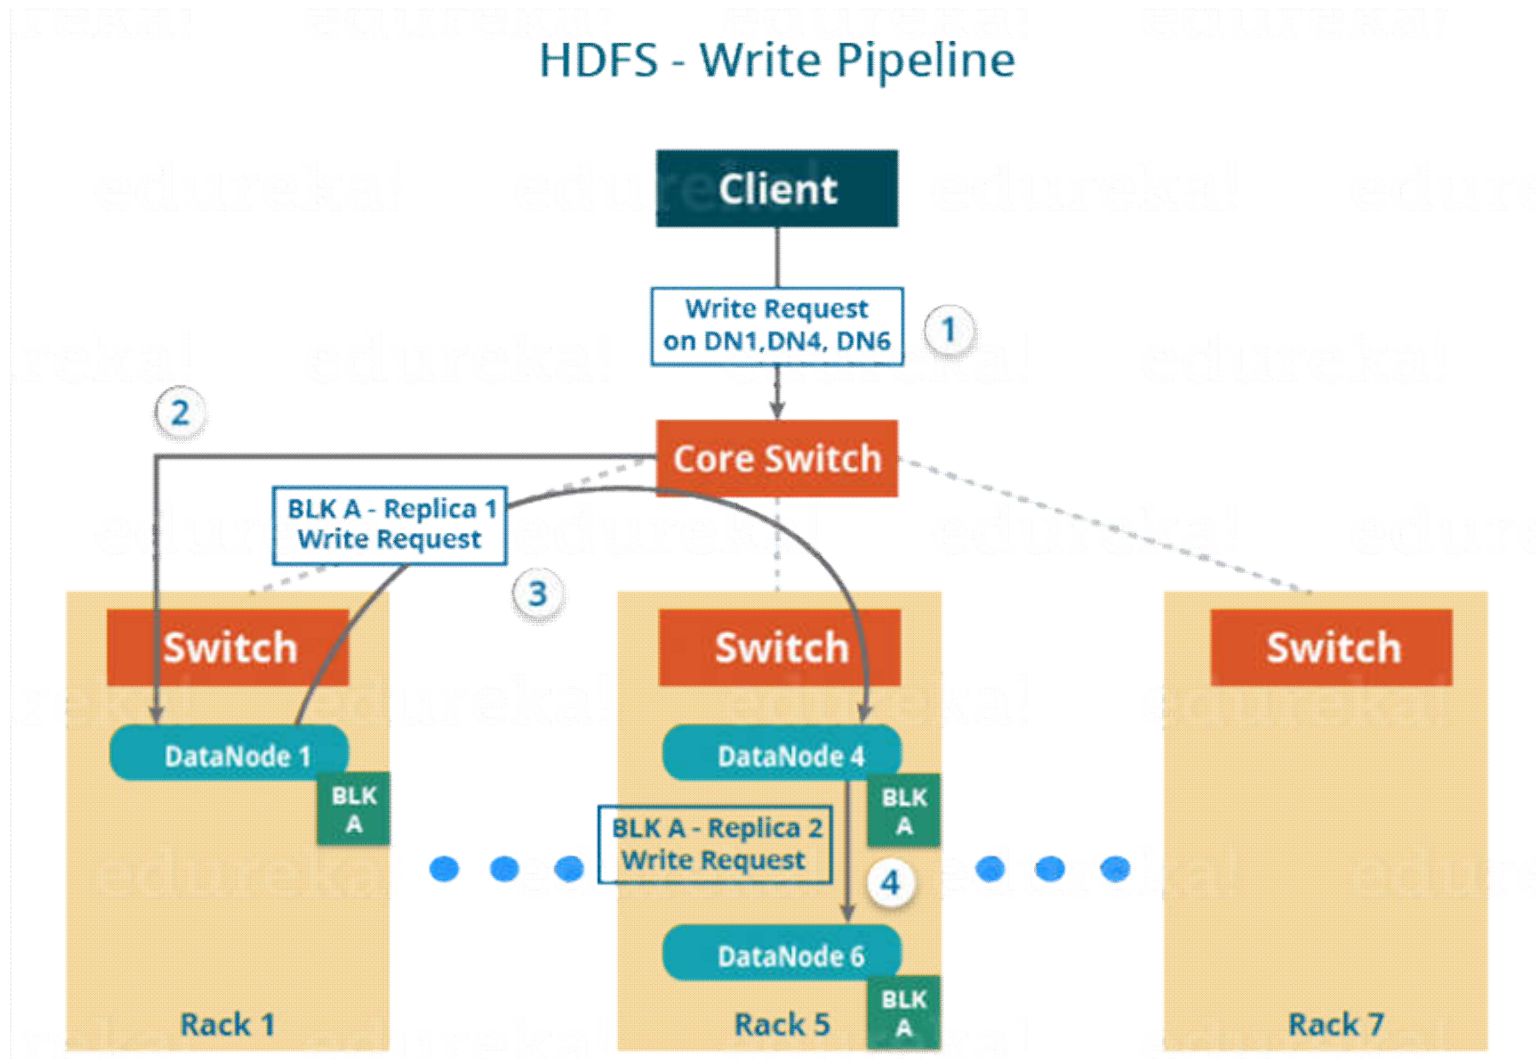
HDFS Write Process :

1). Set up of Pipeline



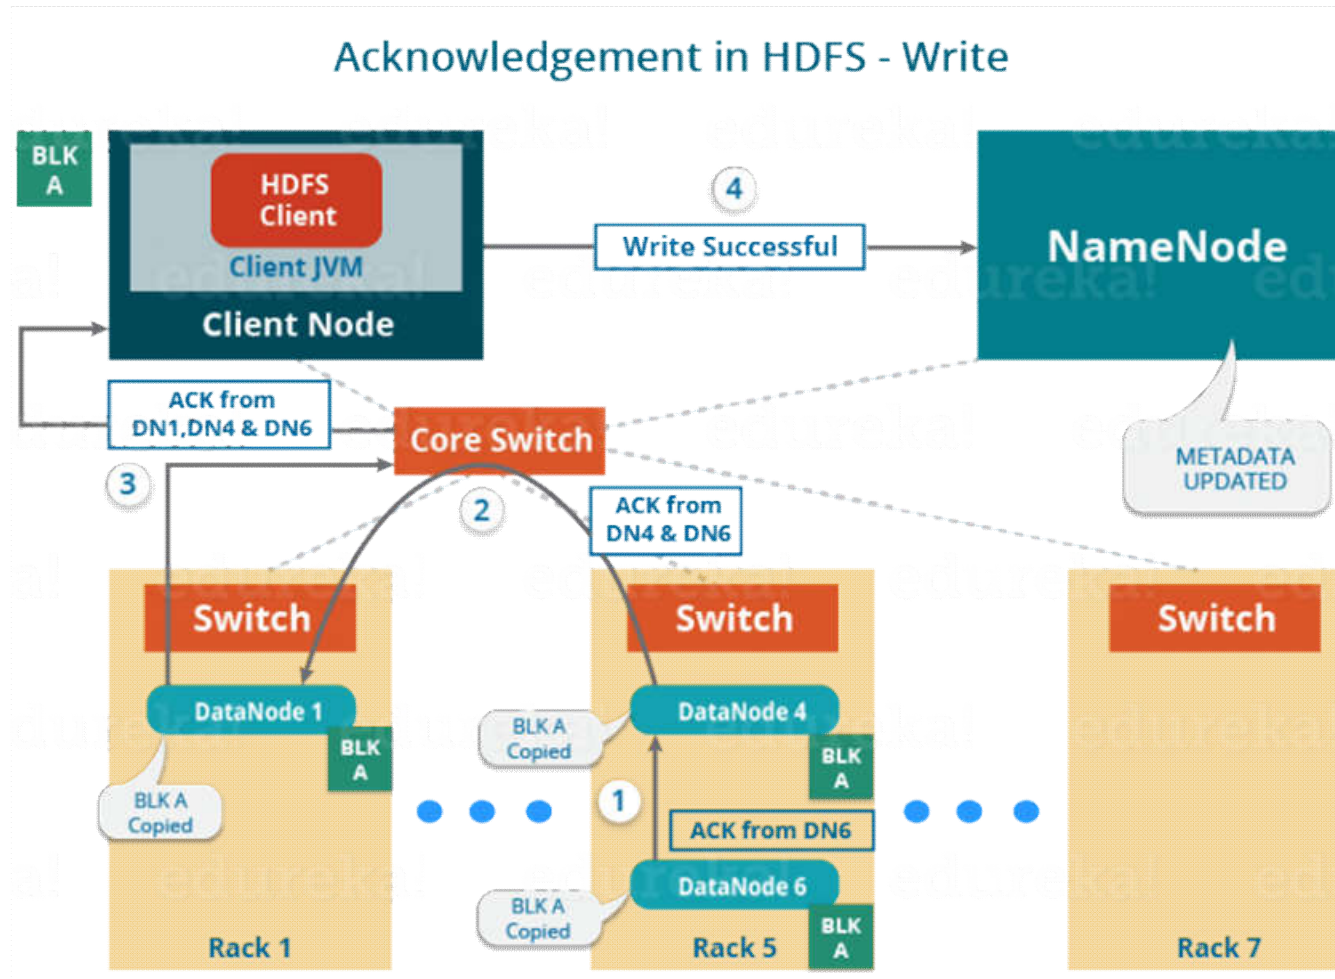
HDFS Write Process :

2). Data streaming and replication

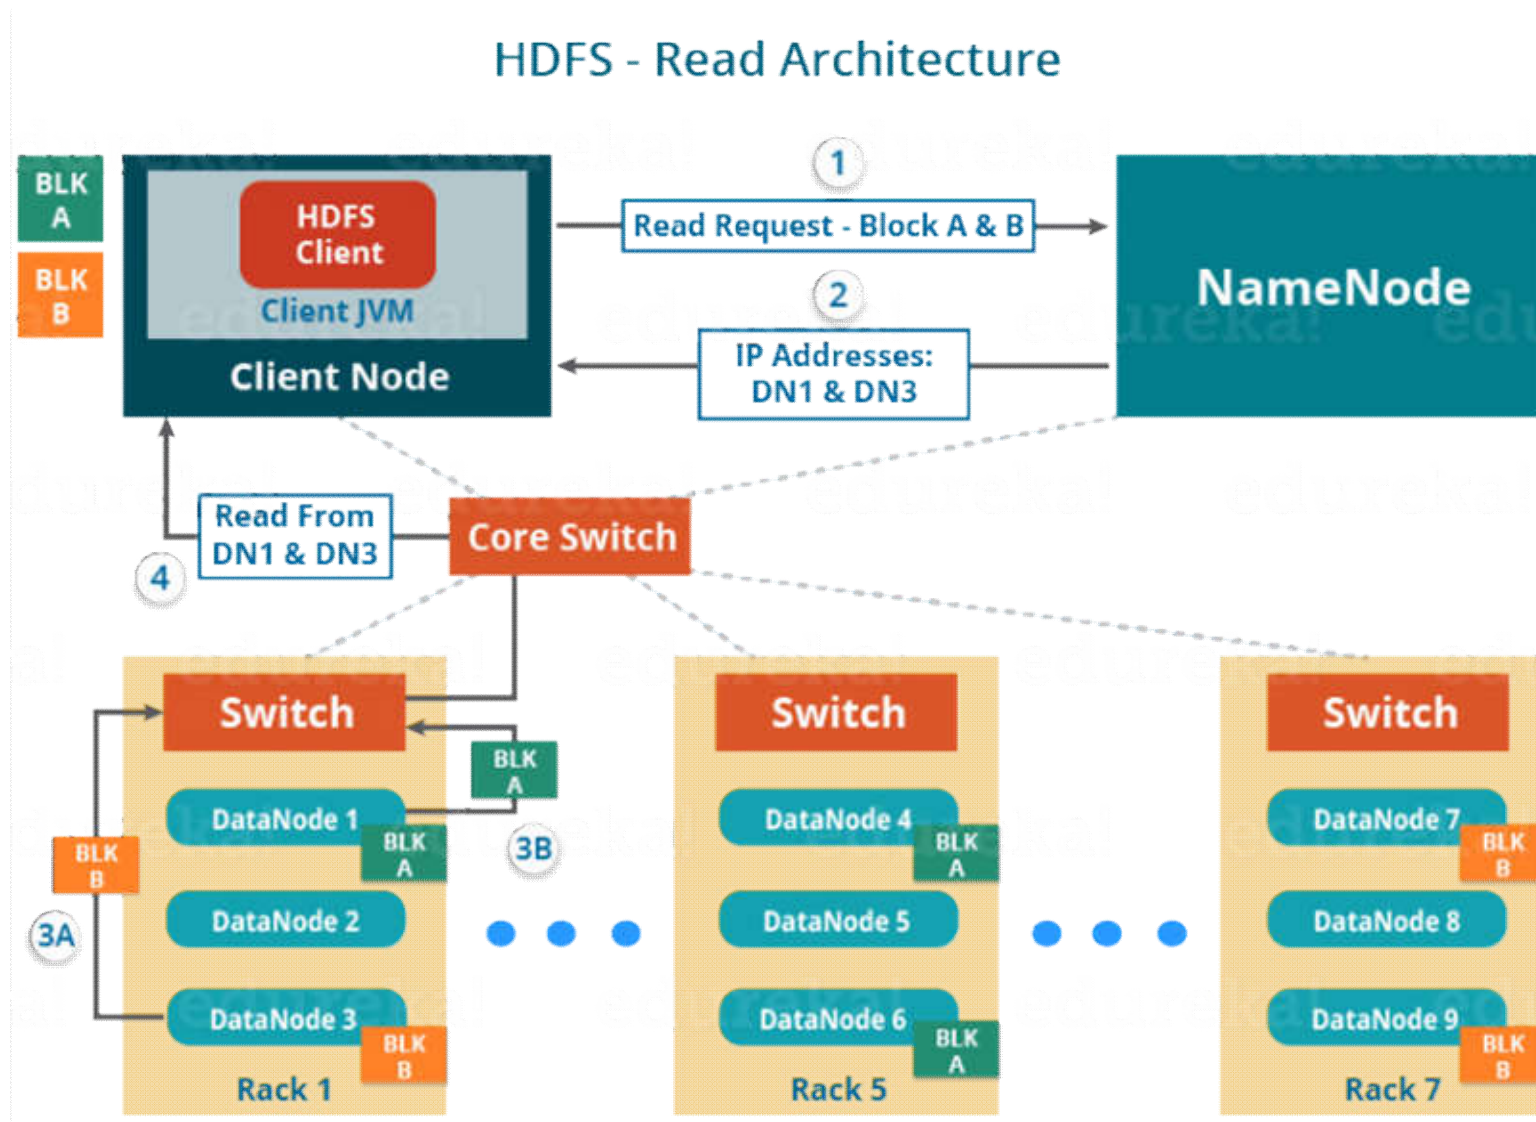


HDFS Write Process :

3). Shutdown of Pipeline



HDFS Read Process :



Hadoop Command :

- File System

Interact to HDFS :

- Interacting with HDFS is primarily performed from the **command line** using the script named **hdfs**. The **hdfs** script has the following usage:

```
$ hadoop fs [-option <arg>]
```

- The **-option** argument is the name of a specific option for the specified command, and **<arg>** is one or more arguments that are specified for this option.
- For example, show help

```
$ hadoop fs -help
```

Interact to HDFS :

- List directory contents

- use `-ls` command: **hdfs dfs -ls**

```
$ hadoop fs -ls
```

- Running the `-ls` command on a new cluster will not return any results. This is because the `-ls` command, without any arguments, will attempt to **display the contents of the user's home directory** on HDFS.
- Providing `-ls` with the forward slash (`/`) as an argument displays the contents of the **root of HDFS**:

```
$ hadoop fs -ls /
```

hdfs dfs -ls

https://www.slideshare.net/eakasit_dpu/introduction-to-hadoop-and-mapreduce-75323926?qid=6dc387d1-f614-4646-a2cb-b4805805235a&v=&b=&from_search=3

Interact to HDFS :

- Creating a directory
 - To create the books directory within HDFS, use the -mkdir command:
hdfs dfs -mkdir
 - For example, create books directory in home directory
- created:

```
$ hadoop fs -mkdir [directory name]
```

```
$ hadoop fs -mkdir books
```

```
$ hadoop fs -ls
```

https://www.slideshare.net/eakasit_dpu/introduction-to-hadoop-and-mapreduce-75323926?qid=6dc387d1-f614-4646-a2cb-b4805805235a&v=&b=&from_search=3

Interact to HDFS :

- Copy Data onto HDFS
 - After a directory has been created for the current user, data can be uploaded to the user's HDFS home directory with the **-put** command:

```
$ hadoop fs -put [source file] [destination file]
```

- For example, copy book file from local to HDFS

```
$ hadoop fs -put pg20417.txt books/pg20417.txt
```

- Use the **-ls** command to verify that pg20417.txt was moved to HDFS:

```
$ hadoop fs -ls books
```

https://www.slideshare.net/eakasit_dpu/introduction-to-hadoop-and-mapreduce-75323926?qid=6dc387d1-f614-4646-a2cb-b4805805235a&v=&b=&from_search=3

Interact to HDFS :

- Retrieve (view) Data from HDFS
 - Multiple commands allow data to be retrieved from HDFS.
 - To simply view the contents of a file, use the **-cat** command. **-cat** reads a file on HDFS and displays its contents to stdout.
 - The following command uses **-cat** to display the contents of `pg20417.txt`

```
$ hadoop fs -cat books/pg20417.txt
```

https://www.slideshare.net/eakasit_dpu/introduction-to-hadoop-and-mapreduce-75323926?qid=6dc387d1-f614-4646-a2cb-b4805805235a&v=&b=&from_search=3

Interact to HDFS :

- Retrieve (view) Data from HDFS
 - Data can also be copied from HDFS to the local filesystem using the **-get** command. The **-get** command is the opposite of the **-put** command:

```
$ hadoop fs -get [source file] [destination file]
```

- For example, This command copies pg20417.txt from HDFS to the local filesystem.

```
$ hadoop fs -get pg20417.txt .
```

https://www.slideshare.net/eakasit_dpu/introduction-to-hadoop-and-mapreduce-75323926?qid=6dc387d1-f614-4646-a2cb-b4805805235a&v=&b=&from_search=3

Wrap up HDFS :



https://www.slideshare.net/eakasit_dpu/introduction-to-hadoop-and-mapreduce-75323926?qid=6dc387d1-f614-4646-a2cb-b4805805235a&v=&b=&from_search=3

Principle of MapReduce :

What is MapReduce :

Google released a paper on MapReduce technology in December, 2004. This became the genesis of the Hadoop Processing Model. So, MapReduce is a programming model that allows us to perform parallel and distributed processing on huge data sets.



Principle of MapReduce :

MapReduce best fit for ? :

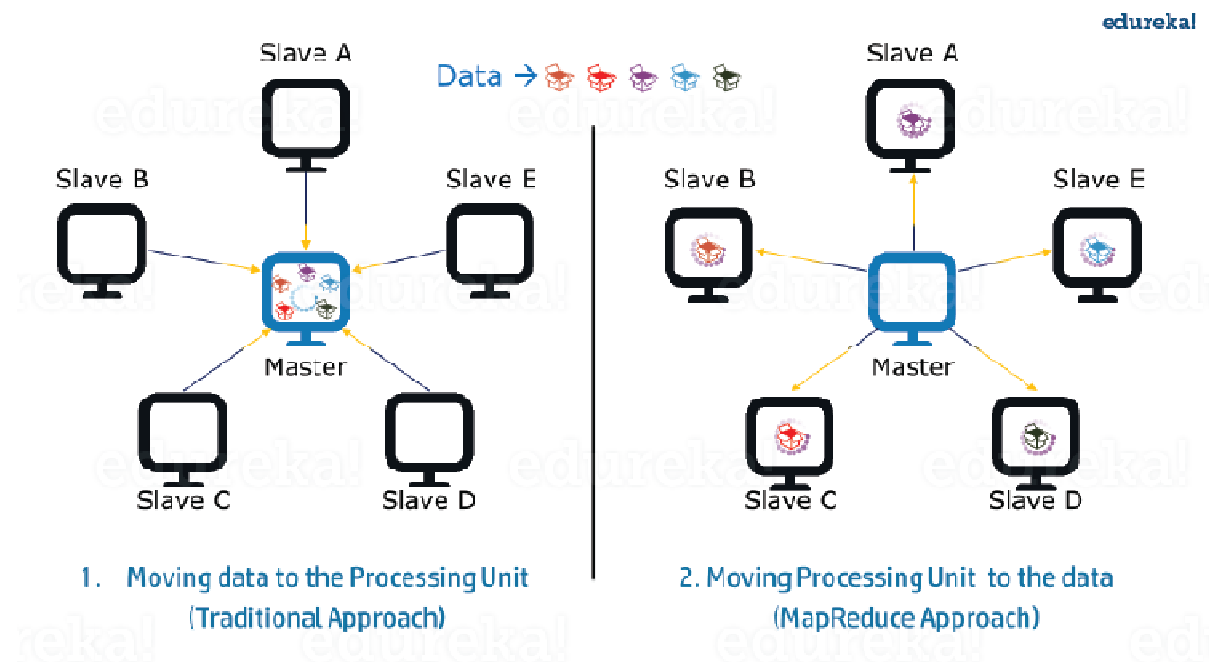
-

Data Processing Model : Traditional vs. Distributed

The biggest advantage of MapReduce :

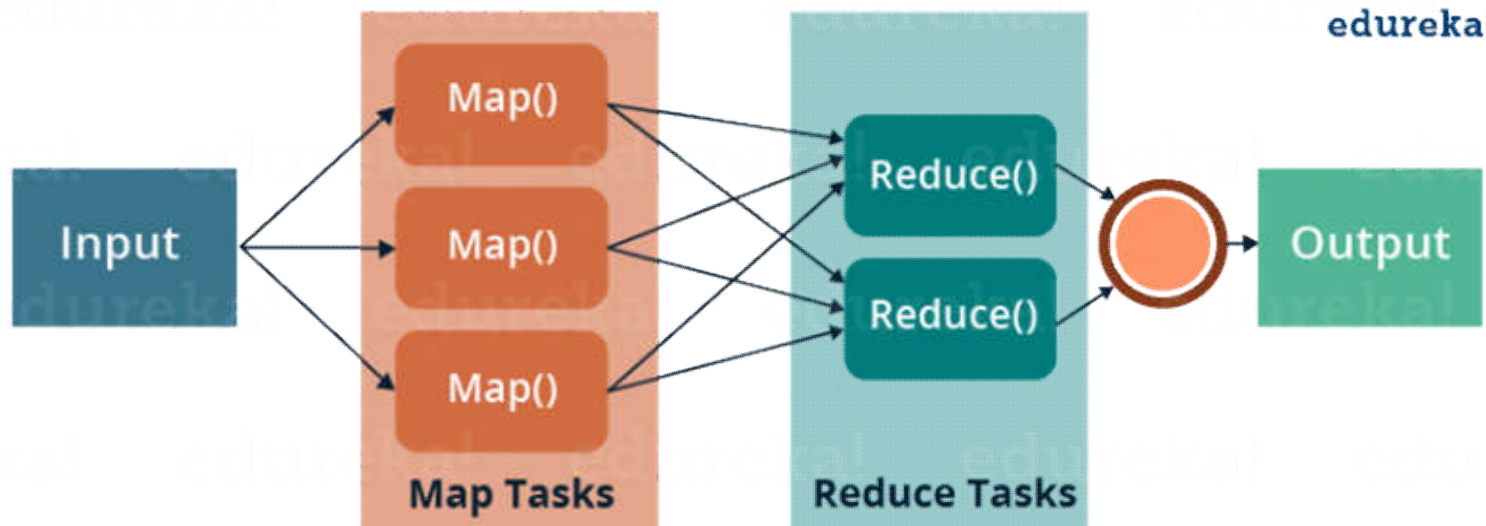
1. Parallel Processing:

2. Data Locality:



MapReduce Concept :

What is MapReduce?

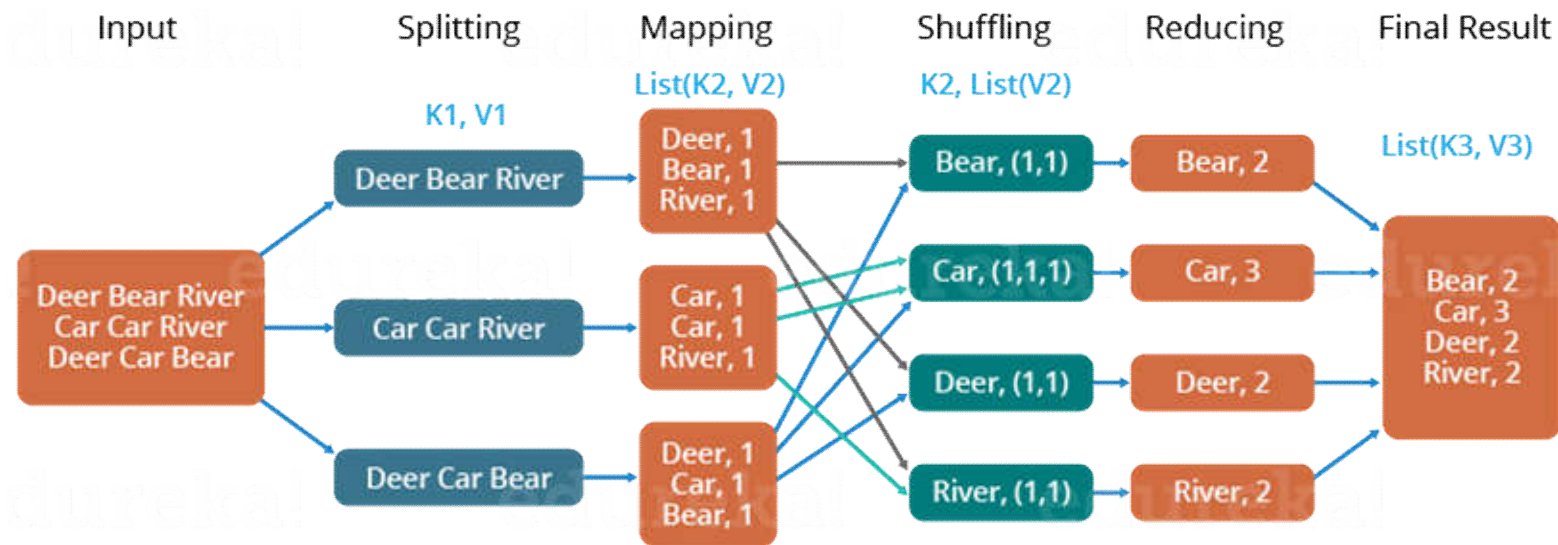


MapReduce Process Flow :



The Overall MapReduce Word Count Process

edureka!



https://www.edureka.co/blog/anatomy-of-a-mapreduce-job-in-apache-hadoop/?utm_source=blog&utm_medium=bottom-related-blog-blogs&utm_campaign=mapreduce-tutorial

MapReduce Programming Model :

```
public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{
        ...
    }

    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        ....
    }

    public static void main(String[] args) throws Exception {

        Configuration conf = new Configuration();

        Job job = Job.getInstance(conf, "word count");

        job.setJarByClass(WordCount.class);

        job.setMapperClass(TokenizerMapper.class);

        job.setCombinerClass(IntSumReducer.class);

        job.setReducerClass(IntSumReducer.class);

        job.setOutputKeyClass(Text.class);

        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

MapReduce Programming Model :

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {

        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}

public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

MapReduce Programming Model :

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

MapReduce Programming Model :

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
  
    Job job = Job.getInstance(conf, "word count");  
  
    job.setJarByClass(WordCount.class);  
  
    job.setMapperClass(TokenizerMapper.class);  
  
    job.setCombinerClass(IntSumReducer.class);  
  
    job.setReducerClass(IntSumReducer.class);  
  
    job.setOutputKeyClass(Text.class);  
  
    job.setOutputValueClass(IntWritable.class);  
  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

MapReduce Programming :

1). Write Code Program

1-1: Create Directory /home/auoychai/units
/home/auoychai/<xxxx>.java

2). Compile & Package

\$javac -classpath hadoop-core-1.2.1.jar -d units <xxxx>.java
\$jar -cvf <xxxx>.jar -C units/ .

3). Run Job-Map/Reduce

Hand-On : File System

1). Create Directory ใน Hadoop

/user/ [lab | dataset]

/lab/[input | output]

/dataset/[csv | txt | json | db

2). Copy file เข้า Hadoop

/home/auoychai/dataset

****** copy file ของแต่ละหมวดหมู่ไปไว้ใน Hadoop Directory /user/dataset/ [...]
ตามกลุ่มของประเภทไฟล์

3). List Directory ใน Hadoop

4). ลบ File / Directory ใน Hadoop

****** Delete file ของแต่ละหมวดหมู่ไปไว้ใน Hadoop Directory /user/dataset/ [...]

Hand-On : File System : MapReduce Programming :

1). Write Code Program

1-1: Create Directory /home/auoychai/units
/home/auoychai/Wordcount.java

2). Compile & Package

\$javac -classpath hadoop-core-1.2.1.jar -d units WordCount.java
\$jar -cvf WordCount.jar -C units/ .

3). Run Job-Map/Reduce

\$hadoop jar WordCount.jar WordCount /user/lab/input/sample.txt
/user/lab/output/o3

Principle HDFS & MapReduce :

Wrap up

The End

Big
data

Shift